

Day49_Project_House_Price_Prediction_Using_All_Regression_Models

July 18, 2025

Housing Price Prediction Using Multiple Regression Algorithms and Pickle

Introduction

Today, we are moving a step ahead from individual regression algorithms and learning how to evaluate and compare multiple regression models in one go using a unified structure.

So far, we have explored:

- Regression models like Linear Regression, Ridge, Lasso, Random Forest, SVM, etc.

In this notebook, we'll:

- Train and evaluate multiple regression algorithms
- Save each trained model as a .pkl file using pickle
- Compare their performance using MAE, MSE, and R^2
- Prepare these models to be used with any frontend like Streamlit

1 Import Required Libraries

```
[1]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import (LinearRegression, Ridge, Lasso, ElasticNet,
    ↳SGDRegressor, HuberRegressor)
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.preprocessing import PolynomialFeatures
from sklearn.pipeline import Pipeline
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
import lightgbm as lgb
import xgboost as xgb
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
import pickle
```

2 Load Dataset

```
[2]: data = pd.read_csv(r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday\work\Github work\ML_Project\USA_Housing.csv")
data.head()
```

```
[2]:
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	\
0	79545.458574	5.682861	7.009188	
1	79248.642455	6.002900	6.730821	
2	61287.067179	5.865890	8.512727	
3	63345.240046	7.188236	5.586729	
4	59982.197226	5.040555	7.839388	

	Avg. Area Number of Bedrooms	Area Population	Price	\
0	4.09	23086.800503	1.059034e+06	
1	3.09	40173.072174	1.505891e+06	
2	5.13	36882.159400	1.058988e+06	
3	3.26	34310.242831	1.260617e+06	
4	4.23	26354.109472	6.309435e+05	

	Address
0	208 Michael Ferry Apt. 674\nLaurabury, NE 3701...
1	188 Johnson Views Suite 079\nLake Kathleen, CA...
2	9127 Elizabeth Stravenue\nDanielstown, WI 06482...
3	USS Barnett\nFPO AP 44820
4	USNS Raymond\nFPO AE 09386

3 Preprocessing

We drop columns that won't help with prediction (Address is non-numeric), and separate the target (Price).

```
[3]: X = data.drop(['Price', 'Address'], axis=1)
y = data['Price']
```

4 Train-Test Split

```
[4]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=0)
```

5 Scale data for models that require it

```
[6]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

6 Define All Regressor Models

```
[7]: models = {
    # Models that need scaled data
    'LinearRegression': LinearRegression(),
    'RobustRegression': HuberRegressor(),
    'RidgeRegression': Ridge(),
    'LassoRegression': Lasso(),
    'ElasticNet': ElasticNet(),
    'PolynomialRegression': Pipeline([
        ('poly', PolynomialFeatures(degree=4)),
        ('linear', LinearRegression())
    ]),
    'SGDRegressor': SGDRegressor(),
    'ANN': MLPRegressor(hidden_layer_sizes=(100,), max_iter=1000),
    'SVM': SVR(),
    'KNN': KNeighborsRegressor(),

    # Models that don't need scaling
    'RandomForest': RandomForestRegressor(),
    'LGBM': lgb.LGBMRegressor(),
    'XGBoost': xgb.XGBRegressor()
}
```

7 Train Models, Evaluate, and Save as .pkl

```
[8]: results = []

for name, model in models.items():
    # Check if model requires scaling
    if name in ['LinearRegression', 'RobustRegression', 'RidgeRegression',
    ↪ 'LassoRegression',
    ↪ 'ElasticNet', 'PolynomialRegression', 'SGDRegressor', 'ANN',
    ↪ 'SVM', 'KNN']:
        model.fit(X_train_scaled, y_train)
        y_pred = model.predict(X_test_scaled)
    else:
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
```

```

# Evaluation
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

results.append({
    'Model': name,
    'MAE': round(mae, 2),
    'MSE': round(mse, 2),
    'R2': round(r2, 4)
})

# Save model
with open(f'{name}.pkl', 'wb') as f:
    pickle.dump(model, f)

print("All models trained, evaluated, and saved as .pkl files.")

```

C:\Users\Lenovo\anaconda3\Lib\site-packages\sklearn\network_multilayer_perceptron.py:780: ConvergenceWarning: Stochastic Optimizer: Maximum iterations (1000) reached and the optimization hasn't converged yet.

warnings.warn(

[LightGBM] [Warning] Found whitespace in feature_names, replace with underlines
 [LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of testing was 0.000492 seconds.

You can set `force_col_wise=true` to remove the overhead.

[LightGBM] [Info] Total Bins 1256

[LightGBM] [Info] Number of data points in the train set: 4000, number of used features: 5

[LightGBM] [Info] Start training from score 1231911.452183

All models trained, evaluated, and saved as .pkl files.

8 Show Model Comparison Results

```

[9]: results_df = pd.DataFrame(results)
results_df.sort_values(by='R2', ascending=False, inplace=True)
results_df.reset_index(drop=True, inplace=True)
print(results_df)

```

	Model	MAE	MSE	R2
0	RobustRegression	82659.92	1.054623e+10	0.9147
1	RidgeRegression	82658.16	1.054893e+10	0.9147
2	SGDRegressor	82567.49	1.054623e+10	0.9147
3	LassoRegression	82657.87	1.054970e+10	0.9146
4	LinearRegression	82657.95	1.054972e+10	0.9146
5	PolynomialRegression	84013.48	1.073798e+10	0.9131

6	LGBM	92133.99	1.309771e+10	0.8940
7	RandomForest	97709.90	1.473850e+10	0.8808
8	XGBoost	101565.19	1.613868e+10	0.8694
9	KNN	105521.78	1.710311e+10	0.8616
10	ElasticNet	121396.83	2.288246e+10	0.8149
11	SVM	282858.36	1.234840e+11	0.0009
12	ANN	1175960.23	1.483829e+12	-11.0052

9 Save Results to CSV

```
[10]: results_df.to_csv('model_evaluation_results.csv', index=False)
      print("Model evaluation saved to model_evaluation_results.csv")
```

Model evaluation saved to model_evaluation_results.csv

10 Final Notes

- All models are saved and can be loaded back in Streamlit or any app using `pickle.load(open('model.pkl', 'rb'))`.
- This helps you pick the best model for production use based on accuracy (R^2), error (MAE), etc.
- You can now build a Streamlit app that allows you to upload input and choose a model to predict!

11 Conclusion

In this project, we developed and compared multiple regression models to predict **housing prices** based on various area-specific features such as average income, house age, number of rooms, number of bedrooms, and area population.

After training and evaluating 13 different models, we found the following insights:

- **Top Performing Models** (based on low MAE, MSE and high R^2 Score):
 - **SGD Regressor**: Lowest MAE (82,567) and high R^2 (0.9147)
 - **Ridge Regression** and **Robust Regression** closely followed with similar performance and strong generalization.
- **Baseline Models** like **Linear Regression** and **Lasso Regression** performed decently with an R^2 of **~0.9146**, suggesting the problem is well-suited for linear approaches.
- **Polynomial Regression** slightly improved performance but at the cost of increased model complexity.
- **Ensemble Models** such as **LGBM** and **Random Forest** performed well but not significantly better than simpler models.
- **XGBoost** and **KNN** underperformed compared to others, with higher error values.

- **SVM** and **ANN** (Artificial Neural Network) showed **very poor performance**, especially the ANN which had an **R^2 of -11.0052**, indicating severe overfitting or improper tuning.

11.1 Final Recommendation

Considering performance and interpretability, **SGD Regressor**, **Ridge Regression**, and **Robust Regression** are the **best choices** for deployment in this use case. These models offer:

- High accuracy
- Low error margins
- Better generalization
- Simpler implementation and tuning

11.2 Future Scope

- Hyperparameter tuning for all models
- More advanced feature engineering or scaling methods
- Experimenting with stacked models or deep learning (with proper normalization and regularization)