

Day82_Deep_Learning_YOLOv8_Object_Detection_Image_Video

September 9, 2025

1 YOLOv8 Object Detection – Complete Beginner Walkthrough

This notebook explains and implements **YOLOv8 (You Only Look Once)** for object detection.

We will cover:

- What is YOLO? Why use it?
- Difference between YOLO and traditional OpenCV Haar Cascades
- Pretrained models & datasets (COCO)
- Installing and setting up Ultralytics
- Running detection on images and videos
- Drawing bounding boxes with OpenCV

2 What is YOLO (You Only Look Once)?

- **YOLO** is a state-of-the-art, real-time object detection algorithm.
- It divides an image into grids and predicts bounding boxes & class probabilities in a **single forward pass**.
- That's why it's **fast and accurate** → suitable for real-time detection.

Use cases: self-driving cars, surveillance, medical imaging, quality inspection in factories, etc.

3 YOLO vs Haar Cascade (OpenCV)

- **Haar Cascades:**
 - Old method in OpenCV for object detection (e.g., face detection).
 - Works by scanning the image with sliding windows & handcrafted features.

- Limitations → only works well for simple tasks like faces; not robust for complex real-world detection.

- **YOLO (Deep Learning-based):**

- Learns automatically from huge datasets (like COCO).
- Handles multiple objects, complex scenes, variations in lighting, orientation, etc.
- Much more accurate & flexible than Haar Cascades.

Conclusion: For modern applications → **YOLO is far better**.

4 Pretrained Models and Datasets

- Ultralytics provides **pretrained YOLO models** on the **COCO dataset** (80 object classes: person, car, dog, etc.).
- Different YOLOv8 model sizes:
 - yolov8n.pt → Nano (fastest, lightest)
 - yolov8s.pt → Small
 - yolov8m.pt → Medium
 - yolov8l.pt → Large
 - yolov8x.pt → Extra large (most accurate but heavy)

You can start with pretrained weights and fine-tune them on your custom dataset later.

5 Setup Environment

This checks:

- Python version
- Torch (PyTorch) version
- GPU/CPU availability

```
[ ]: # Install YOLO library
!pip install ultralytics
```

```
[1]: # Import and check environment
import ultralytics
ultralytics.checks()
```

Ultralytics 8.3.196 Python-3.13.5 torch-2.8.0+cpu CPU (Intel Core i3-5005U
2.00GHz)
Setup complete (4 CPUs, 7.9 GB RAM, 139.2/237.8 GB disk)

6 Load a Pretrained YOLO Model

```
[2]: from ultralytics import YOLO
import numpy as np

# Load a pretrained YOLOv8 Nano model
model = YOLO("yolov8n.pt", "v8")
```

7 Run Prediction on an Image

- conf=0.25 → Minimum confidence threshold
- save=True → Saves results with bounding boxes to runs/detect folder

```
[3]: detection_output = model.predict(
    source=r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github_
work\Computer_Vision\YOLO\Test_Image_and_Video\Image.JPG",
    conf=0.25,
    save=True
)

# Display raw tensor output
print(detection_output)

# Convert first detection result to NumPy
print(detection_output[0].numpy())
```

image 1/1 C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github
work\Computer_Vision\YOLO\Test_Image_and_Video\Image.JPG: 416x640 30 persons, 10
cars, 1 bus, 3 trucks, 4 traffic lights, 1061.3ms
Speed: 55.6ms preprocess, 1061.3ms inference, 26.7ms postprocess per image at
shape (1, 3, 416, 640)
Results saved to C:\Users\Lenovo\OneDrive\Desktop\Python Everyday

work\Github work\Computer_Vision\YOLO\runs\detect\predict
[ultralytics.engine.results.Results object with attributes:

boxes: ultralytics.engine.results.Boxes object

keypoints: None

masks: None

names: {0: 'person', 1: 'bicycle', 2: 'car', 3: 'motorcycle', 4: 'airplane', 5:
'bus', 6: 'train', 7: 'truck', 8: 'boat', 9: 'traffic light', 10: 'fire
hydrant', 11: 'stop sign', 12: 'parking meter', 13: 'bench', 14: 'bird', 15:

'cat', 16: 'dog', 17: 'horse', 18: 'sheep', 19: 'cow', 20: 'elephant', 21: 'bear', 22: 'zebra', 23: 'giraffe', 24: 'backpack', 25: 'umbrella', 26: 'handbag', 27: 'tie', 28: 'suitcase', 29: 'frisbee', 30: 'skis', 31: 'snowboard', 32: 'sports ball', 33: 'kite', 34: 'baseball bat', 35: 'baseball glove', 36: 'skateboard', 37: 'surfboard', 38: 'tennis racket', 39: 'bottle', 40: 'wine glass', 41: 'cup', 42: 'fork', 43: 'knife', 44: 'spoon', 45: 'bowl', 46: 'banana', 47: 'apple', 48: 'sandwich', 49: 'orange', 50: 'broccoli', 51: 'carrot', 52: 'hot dog', 53: 'pizza', 54: 'donut', 55: 'cake', 56: 'chair', 57: 'couch', 58: 'potted plant', 59: 'bed', 60: 'dining table', 61: 'toilet', 62: 'tv', 63: 'laptop', 64: 'mouse', 65: 'remote', 66: 'keyboard', 67: 'cell phone', 68: 'microwave', 69: 'oven', 70: 'toaster', 71: 'sink', 72: 'refrigerator', 73: 'book', 74: 'clock', 75: 'vase', 76: 'scissors', 77: 'teddy bear', 78: 'hair drier', 79: 'toothbrush'}

obb: None

```
orig_img: array([[[ 27,  28,  26],
                  [ 25,  26,  24],
                  [ 26,  27,  25],
                  ...,
                  [ 63,  72,  69],
                  [ 29,  37,  37],
                  [ 18,  28,  28]],

                [[ 28,  29,  27],
                  [ 25,  26,  24],
                  [ 27,  28,  26],
                  ...,
                  [ 43,  50,  47],
                  [ 11,  19,  18],
                  [ 44,  55,  53]],

                [[ 29,  30,  28],
                  [ 27,  28,  26],
                  [ 28,  29,  27],
                  ...,
                  [ 31,  38,  33],
                  [ 19,  29,  23],
                  [ 83,  92,  89]],

                ...,

                [[ 94,  83,  69],
                  [ 94,  83,  69],
                  [ 92,  81,  67],
                  ...,
                  [113, 107, 100],
                  [107, 101,  94],
                  [103,  97,  90]],
```

```

[[ 89, 78, 64],
 [ 88, 77, 63],
 [ 85, 74, 60],
 ...,
 [129, 123, 116],
 [129, 123, 116],
 [133, 127, 120]],

[[ 83, 72, 58],
 [ 81, 70, 56],
 [ 77, 66, 52],
 ...,
 [132, 126, 119],
 [129, 123, 116],
 [138, 132, 125]]], dtype=uint8)
orig_shape: (582, 910)
path: 'C:\\Users\\Lenovo\\OneDrive\\Desktop\\Python Everyday work\\Github
work\\Computer_Vision\\YOLO\\Test_Image_and_Video\\Image.JPG'
probs: None
save_dir: 'C:\\Users\\Lenovo\\OneDrive\\Desktop\\Python Everyday work\\Github
work\\Computer_Vision\\YOLO\\runs\\detect\\predict'
speed: {'preprocess': 55.56060001254082, 'inference': 1061.301099951379,
'postprocess': 26.681900024414062}]
ultralalytics.engine.results.Results object with attributes:

boxes: ultralytics.engine.results.Boxes object
keypoints: None
masks: None
names: {0: 'person', 1: 'bicycle', 2: 'car', 3: 'motorcycle', 4: 'airplane', 5:
'bus', 6: 'train', 7: 'truck', 8: 'boat', 9: 'traffic light', 10: 'fire
hydrant', 11: 'stop sign', 12: 'parking meter', 13: 'bench', 14: 'bird', 15:
'cat', 16: 'dog', 17: 'horse', 18: 'sheep', 19: 'cow', 20: 'elephant', 21:
'bear', 22: 'zebra', 23: 'giraffe', 24: 'backpack', 25: 'umbrella', 26:
'handbag', 27: 'tie', 28: 'suitcase', 29: 'frisbee', 30: 'skis', 31:
'snowboard', 32: 'sports ball', 33: 'kite', 34: 'baseball bat', 35: 'baseball
glove', 36: 'skateboard', 37: 'surfboard', 38: 'tennis racket', 39: 'bottle',
40: 'wine glass', 41: 'cup', 42: 'fork', 43: 'knife', 44: 'spoon', 45: 'bowl',
46: 'banana', 47: 'apple', 48: 'sandwich', 49: 'orange', 50: 'broccoli', 51:
'carrot', 52: 'hot dog', 53: 'pizza', 54: 'donut', 55: 'cake', 56: 'chair', 57:
'couch', 58: 'potted plant', 59: 'bed', 60: 'dining table', 61: 'toilet', 62:
'tv', 63: 'laptop', 64: 'mouse', 65: 'remote', 66: 'keyboard', 67: 'cell phone',
68: 'microwave', 69: 'oven', 70: 'toaster', 71: 'sink', 72: 'refrigerator', 73:
'book', 74: 'clock', 75: 'vase', 76: 'scissors', 77: 'teddy bear', 78: 'hair
drier', 79: 'toothbrush'}
obb: None
orig_img: array([[[ 27, 28, 26],
 [ 25, 26, 24],
 [ 26, 27, 25],

```

```

...,
[ 63, 72, 69],
[ 29, 37, 37],
[ 18, 28, 28]],

[[ 28, 29, 27],
[ 25, 26, 24],
[ 27, 28, 26],
...,
[ 43, 50, 47],
[ 11, 19, 18],
[ 44, 55, 53]],

[[ 29, 30, 28],
[ 27, 28, 26],
[ 28, 29, 27],
...,
[ 31, 38, 33],
[ 19, 29, 23],
[ 83, 92, 89]],

...,

[[ 94, 83, 69],
[ 94, 83, 69],
[ 92, 81, 67],
...,
[113, 107, 100],
[107, 101, 94],
[103, 97, 90]],

[[ 89, 78, 64],
[ 88, 77, 63],
[ 85, 74, 60],
...,
[129, 123, 116],
[129, 123, 116],
[133, 127, 120]],

[[ 83, 72, 58],
[ 81, 70, 56],
[ 77, 66, 52],
...,
[132, 126, 119],
[129, 123, 116],
[138, 132, 125]]], dtype=uint8)
orig_shape: (582, 910)
path: 'C:\\Users\\Lenovo\\OneDrive\\Desktop\\Python Everyday work\\Github

```

```
work\\Computer_Vision\\YOLO\\Test_Image_and_Video\\Image.JPG'
probs: None
save_dir: None
speed: {'preprocess': 55.56060001254082, 'inference': 1061.301099951379,
'postprocess': 26.681900024414062}
```

7.1 YOLOv8 Detection Output (Image Detection)

This is the detection result for Uploaded Image:



8 Load Class Labels and Colors

```
[4]: import random

# Load COCO class names
with open(r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github_\
work\Computer_Vision\YOLO\coco.txt", "r") as f:
    class_list = f.read().split("\n")

# Assign random colors for bounding boxes
detection_colors = [(random.randint(0,255), random.randint(0,255), random.
    randint(0,255)) for _ in class_list]

# Run YOLO detection on a different image
results = model.predict(
```

```

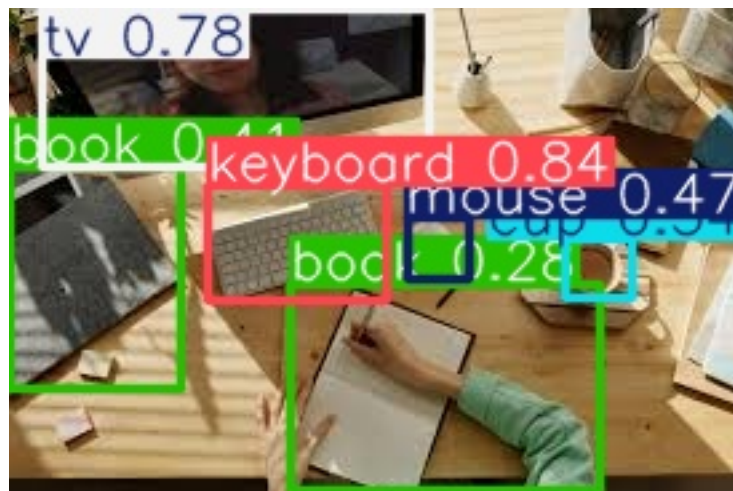
source=r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github_
work\Computer_Vision\YOLO\Test_Image_and_Video\Image_1.JPG",
conf=0.25,
save=True
)

```

image 1/1 C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github work\Computer_Vision\YOLO\Test_Image_and_Video\Image_1.JPG: 448x640 1 cup, 1 tv, 1 mouse, 1 keyboard, 2 books, 681.0ms
Speed: 14.7ms preprocess, 681.0ms inference, 12.0ms postprocess per image at shape (1, 3, 448, 640)
Results saved to C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github work\Computer_Vision\YOLO\runs\detect\predict

8.1 YOLOv8 Detection Output (Image Detection)

This is the detection result for **Any Image** using COCO class labels:



9 Object Detection on Video

```

[11]: import cv2
import random
from ultralytics import YOLO

# Load YOLO model
model = YOLO("yolov8n.pt", "v8")

# Load COCO classes
with open(r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday work\Github_
work\Computer_Vision\YOLO\coco.txt", "r", encoding="utf-8") as f:
    class_list = f.read().splitlines()

```



```

# Assign random unique colors for each class
detection_colors = [(random.randint(0,255), random.randint(0,255), random.
↳ randint(0,255)) for _ in class_list]

# Open a video file
cap = cv2.VideoCapture(r"C:\Users\Lenovo\OneDrive\Desktop\Python Everyday_
↳ work\Github work\Computer_Vision\YOLO\Test_Image_and_Video\Video.mp4")

if not cap.isOpened():
    print("Cannot open video")
    exit()

# Create a resizable window with minimize/maximize/close buttons
cv2.namedWindow("YOLOv8 Object Detection", cv2.WINDOW_NORMAL)

# Optionally maximize window (depends on OS support)
cv2.setWindowProperty("YOLOv8 Object Detection", cv2.WND_PROP_AUTOSIZE, cv2.
↳ WINDOW_NORMAL)

while True:
    ret, frame = cap.read()
    if not ret:
        print("End of stream")
        break

    # Run YOLO prediction
    results = model.predict(source=[frame], conf=0.45, save=False)
    detections = results[0].numpy()

    if len(detections) != 0:
        for i in range(len(results[0])):
            boxes = results[0].boxes
            box = boxes[i]
            clsID = int(box.cls.numpy()[0])
            conf = box.conf.numpy()[0]
            bb = box.xyxy.numpy()[0]

            # Different color for each object class
            color = detection_colors[clsID]

            # Draw bounding box
            cv2.rectangle(frame, (int(bb[0]), int(bb[1])), (int(bb[2]),
↳ int(bb[3])), color, 2)

            # Put class name + confidence
            cv2.putText(frame, f"{class_list[clsID]} {round(conf,2)}",

```

```

        (int(bb[0]), int(bb[1]) - 10),
        cv2.FONT_HERSHEY_SIMPLEX, 0.7, color, 2, cv2.LINE_AA)

    # Show video in a resizable window (with minimize/maximize/close)
    cv2.imshow("YOLOv8 Object Detection", frame)

    # Exit on ESC key
    if cv2.waitKey(1) == 27:
        break

cap.release()
cv2.destroyAllWindows()

```

0: 384x640 5 persons, 2 chairs, 1 laptop, 588.6ms
 Speed: 17.5ms preprocess, 588.6ms inference, 13.6ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 1 laptop, 511.7ms
 Speed: 15.8ms preprocess, 511.7ms inference, 13.9ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 535.4ms
 Speed: 12.5ms preprocess, 535.4ms inference, 13.2ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 1 laptop, 611.9ms
 Speed: 17.4ms preprocess, 611.9ms inference, 11.3ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 4 laptops, 585.3ms
 Speed: 18.8ms preprocess, 585.3ms inference, 12.6ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 1 laptop, 461.9ms
 Speed: 16.3ms preprocess, 461.9ms inference, 11.9ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 526.2ms
 Speed: 18.5ms preprocess, 526.2ms inference, 14.3ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 529.5ms
 Speed: 20.3ms preprocess, 529.5ms inference, 13.7ms postprocess per image at
 shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 602.9ms
 Speed: 17.6ms preprocess, 602.9ms inference, 14.7ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 643.2ms

Speed: 31.6ms preprocess, 643.2ms inference, 14.0ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 391.1ms

Speed: 18.1ms preprocess, 391.1ms inference, 11.6ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 397.3ms

Speed: 17.3ms preprocess, 397.3ms inference, 12.9ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 452.6ms

Speed: 17.5ms preprocess, 452.6ms inference, 13.9ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 385.9ms

Speed: 17.8ms preprocess, 385.9ms inference, 41.4ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 413.7ms

Speed: 15.0ms preprocess, 413.7ms inference, 16.6ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 3 laptops, 861.1ms

Speed: 22.4ms preprocess, 861.1ms inference, 15.7ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 4 laptops, 687.8ms

Speed: 19.8ms preprocess, 687.8ms inference, 19.3ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 1607.5ms

Speed: 21.6ms preprocess, 1607.5ms inference, 69.8ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 725.1ms

Speed: 24.1ms preprocess, 725.1ms inference, 14.4ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 4 laptops, 737.5ms

Speed: 19.9ms preprocess, 737.5ms inference, 15.1ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 662.4ms

Speed: 21.4ms preprocess, 662.4ms inference, 19.4ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 859.1ms
Speed: 23.0ms preprocess, 859.1ms inference, 16.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 877.1ms
Speed: 17.6ms preprocess, 877.1ms inference, 15.8ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 1069.3ms
Speed: 15.8ms preprocess, 1069.3ms inference, 74.0ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 652.0ms
Speed: 29.3ms preprocess, 652.0ms inference, 14.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 765.5ms
Speed: 42.4ms preprocess, 765.5ms inference, 15.5ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 450.8ms
Speed: 17.2ms preprocess, 450.8ms inference, 13.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 494.1ms
Speed: 19.7ms preprocess, 494.1ms inference, 15.5ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 471.0ms
Speed: 15.5ms preprocess, 471.0ms inference, 17.3ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 10 laptops, 1067.6ms
Speed: 21.2ms preprocess, 1067.6ms inference, 12.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 463.6ms
Speed: 15.7ms preprocess, 463.6ms inference, 15.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 7 laptops, 547.5ms
Speed: 15.5ms preprocess, 547.5ms inference, 13.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 4 laptops, 527.5ms
Speed: 16.9ms preprocess, 527.5ms inference, 13.2ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 6 laptops, 655.8ms
Speed: 28.1ms preprocess, 655.8ms inference, 20.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 464.6ms
Speed: 19.0ms preprocess, 464.6ms inference, 12.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 2 laptops, 400.9ms
Speed: 16.4ms preprocess, 400.9ms inference, 12.7ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 4 laptops, 404.7ms
Speed: 15.9ms preprocess, 404.7ms inference, 12.4ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 416.9ms
Speed: 15.5ms preprocess, 416.9ms inference, 12.8ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 411.1ms
Speed: 14.7ms preprocess, 411.1ms inference, 12.7ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 518.7ms
Speed: 14.6ms preprocess, 518.7ms inference, 17.0ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 1179.2ms
Speed: 15.2ms preprocess, 1179.2ms inference, 15.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 577.5ms
Speed: 21.1ms preprocess, 577.5ms inference, 15.0ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 5 laptops, 777.0ms
Speed: 17.6ms preprocess, 777.0ms inference, 13.3ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 6 laptops, 765.2ms
Speed: 29.7ms preprocess, 765.2ms inference, 12.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 4 laptops, 595.3ms
Speed: 18.1ms preprocess, 595.3ms inference, 12.4ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 5 laptops, 766.1ms
Speed: 19.3ms preprocess, 766.1ms inference, 13.3ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 4 laptops, 616.9ms
Speed: 26.8ms preprocess, 616.9ms inference, 15.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 4 laptops, 651.6ms
Speed: 16.0ms preprocess, 651.6ms inference, 18.3ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 465.1ms
Speed: 15.9ms preprocess, 465.1ms inference, 18.3ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 684.0ms
Speed: 17.7ms preprocess, 684.0ms inference, 14.7ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 5 laptops, 467.4ms
Speed: 17.2ms preprocess, 467.4ms inference, 14.7ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 645.4ms
Speed: 17.5ms preprocess, 645.4ms inference, 19.2ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 715.2ms
Speed: 30.9ms preprocess, 715.2ms inference, 20.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 459.1ms
Speed: 18.4ms preprocess, 459.1ms inference, 22.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 511.7ms
Speed: 19.3ms preprocess, 511.7ms inference, 16.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 620.6ms
Speed: 19.0ms preprocess, 620.6ms inference, 15.5ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 450.2ms
Speed: 17.2ms preprocess, 450.2ms inference, 17.5ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 466.1ms
Speed: 15.4ms preprocess, 466.1ms inference, 19.5ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 560.2ms
Speed: 15.4ms preprocess, 560.2ms inference, 13.4ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 423.8ms
Speed: 22.3ms preprocess, 423.8ms inference, 13.6ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 673.3ms
Speed: 15.3ms preprocess, 673.3ms inference, 12.8ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 410.1ms
Speed: 17.8ms preprocess, 410.1ms inference, 11.8ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 489.0ms
Speed: 32.6ms preprocess, 489.0ms inference, 13.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 411.0ms
Speed: 15.5ms preprocess, 411.0ms inference, 12.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 1 laptop, 407.0ms
Speed: 17.7ms preprocess, 407.0ms inference, 13.9ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 417.6ms
Speed: 18.4ms preprocess, 417.6ms inference, 12.0ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 475.4ms
Speed: 17.4ms preprocess, 475.4ms inference, 13.4ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 404.3ms
Speed: 17.1ms preprocess, 404.3ms inference, 13.1ms postprocess per image at
shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 4 laptops, 432.2ms
Speed: 15.7ms preprocess, 432.2ms inference, 11.4ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 3 laptops, 401.1ms

Speed: 17.2ms preprocess, 401.1ms inference, 13.7ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 3 persons, 2 chairs, 3 laptops, 404.4ms

Speed: 15.5ms preprocess, 404.4ms inference, 13.9ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 1 cup, 2 chairs, 2 laptops, 449.3ms

Speed: 18.8ms preprocess, 449.3ms inference, 15.0ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 421.0ms

Speed: 17.4ms preprocess, 421.0ms inference, 14.5ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 435.4ms

Speed: 18.0ms preprocess, 435.4ms inference, 11.3ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 3 laptops, 424.9ms

Speed: 14.6ms preprocess, 424.9ms inference, 13.0ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 2 laptops, 610.1ms

Speed: 21.0ms preprocess, 610.1ms inference, 9.6ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 4 persons, 2 chairs, 3 laptops, 423.7ms

Speed: 18.0ms preprocess, 423.7ms inference, 14.6ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 8 persons, 2 chairs, 3 laptops, 427.4ms

Speed: 14.6ms preprocess, 427.4ms inference, 14.4ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 11 persons, 2 chairs, 3 laptops, 409.6ms

Speed: 17.1ms preprocess, 409.6ms inference, 13.9ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 8 persons, 2 chairs, 2 laptops, 405.1ms

Speed: 15.5ms preprocess, 405.1ms inference, 12.7ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 2 laptops, 555.5ms

Speed: 22.3ms preprocess, 555.5ms inference, 16.8ms postprocess per image at

shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 2 laptops, 435.4ms

Speed: 17.5ms preprocess, 435.4ms inference, 13.5ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 2 laptops, 1388.2ms

Speed: 56.8ms preprocess, 1388.2ms inference, 14.1ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 2 laptops, 1372.2ms


Speed: 116.9ms preprocess, 1372.2ms inference, 16.4ms postprocess per image at shape (1, 3, 384, 640)

0: 384x640 5 persons, 2 chairs, 2 laptops, 454.6ms

Speed: 20.7ms preprocess, 454.6ms inference, 14.5ms postprocess per image at shape (1, 3, 384, 640)

9.1 YOLOv8 Detection Output (Video Detection)

This is the detection result for the video (screenshot saved as **Video.png**):



10 Conclusion – Day82 YOLOv8 Project

In this notebook, we successfully:

- Installed and set up Ultralytics YOLOv8

- Compared YOLO with Haar Cascades (OpenCV)
- Used **pretrained models on COCO dataset**
- Detected objects in **images** and **videos**
- Displayed results directly inside the notebook

10.1 Key Insights:

- YOLOv8 is **faster and more accurate** than traditional methods like Haar Cascades.
- Pretrained models on COCO allow instant use for 80+ classes.
- With simple code, YOLOv8 can handle both **images** and **video streams**.

10.2 Next Steps:

- Try YOLOv8 with **live webcam detection**
- Fine-tune YOLOv8 on a **custom dataset** for your own project
- Explore advanced tasks like **segmentation** and **pose estimation**

Thanks for following along!